BUSINESS REPORT

ALPHACOM TELECOM

CUSTOMER CHURN PREDICTION

PREPARED BY-
AYUSHI NIGAM

PGP-DS
MAY'25

Table of Contents

Introduction 6
Objective 7
Need Of Studybor 8
Problem Statement 9
Solutions 10
- Data Overview 10
- Head of Data 13
- Tail of Data 14
- Descriptive Statistics 15

Data Preprocessing 16
- Duplicate Check on Data 16
- Missing Value Check on Data & Data Standardization 16
- Fixing Categorical Inconsistencies and Data Validation 18
- Outlier Detection 18

Exploratory Data Analysis 21
- Univariate Analysis 21
- Bivariate Analysis 27
- Multi-Variate Analysis 32

Basic Model Building & Evaluation 36
- Metric of Choice 36
- Data Preparation 36
- Handling Class Imbalance 37

Baseline Model 37
Model Training Approach 38
Model Performance 39
Performance Interpretation 40
Model Selection Rationale 41
Model Performance Visualization-ROC Curve Analysis 42
Advance Model Building & Evaluation 43
Models Implemented 43
Model Evaluation Metrics 46
Performance Analysis Chart 47
Hyperparameter Tuning 48
Model Performance Comparison and Final Model Selection 50
Final Model Evaluation 51
Actionable Insights & Recommendations 55
Key Insights 55
Business Recommendations 56

Table of Figure

Fig. 1.Info of Data 12
Fig. 2.Head of Data 13
Fig. 3.Tail of Data 14
Fig. 4.Descriptive Analysis-Numerical 15
Fig. 5.Descriptive Analysis-Categorical 15
Fig. 6.Duplicate Check 16
Fig. 7.Missing Values 17
Fig. 8.Outlier Detection 19
Fig. 9.Total Outliers 20
Fig. 10.Univariate- Distribution Plots 22
Fig. 11.Univariate-Categorical 24
Fig. 12.Univariate-Churn Distribution 25
Fig. 13.Univariate- Charges Density 25
Fig. 14.Univariate- Customer Distribution (Churn Status, Contract Type, and Internet Service) 26
Fig. 15.Bivariate Analysis of Customer Metrics vs Churn (Tenure, Monthly Charges, Total Charges) 27
Fig. 16.Bivariate Analysis of Contract Type vs Customer Churn 28
Fig. 17.Bivariate- Relationship between Tenure and Total Charges by Churn Status 29
Fig. 18.Bivariate-Monthly Charges by Churn 30
Fig. 19.Bivariate- Tenure and Contract Type vs Customer Churn 31
Fig. 20.Bivariate-Churn Rate across Contract Types 32
Fig. 21.Correlation Analysis of Customer Metrics and Churn 33

Fig. 22.Multivariate-Relationship Between Numerical Variables and Churn 34
Fig. 23. Class Imbalance 37
Fig. 24.Model Performance 39
Fig. 25.Baseline Model Confusion Matrix 40
Fig. 26. Model Performance Visualization-ROC Curve 42
Fig. 27.Model Performance Evaluation 46
Fig. 28.Model Performance Chart 48
Fig. 29. Hyperparameter Tuning Results 49
Fig. 30.Final Model-Performance Metrics 51
Fig. 31.Classification Matrix 51
Fig. 32.Confusion Matrix 53
Fig. 33.ROC Curve 54

Introduction

In today's highly competitive telecommunications industry, customer retention has become as critical as customer acquisition. AlphaCom, a leading telecom provider, is currently facing a significant increase in customer churn, which poses a direct threat to its revenue streams and market positioning. Despite offering a diverse portfolio of services and competitive pricing, the company is unable to effectively curb customer attrition.

Customer churn is a complex phenomenon influenced by a combination of demographic, behavioral, and service-related factors such as contract type, service usage, billing preferences, and tenure. Traditional reactive approaches to retention have proven insufficient, as they fail to identify at-risk customers before they decide to leave.

This study focuses on leveraging data-driven techniques to better understand churn behavior. By analyzing historical customer data and identifying patterns associated with churn, the organization can transition from reactive problem-solving to proactive decision-making, thereby enhancing customer retention strategies and long-term profitability.

Objective

The primary objective of this study is to develop a robust predictive model that can accurately identify customers who are at high risk of churning. By leveraging historical customer data, the model aims to classify customers based on their likelihood of leaving the company.

In addition to prediction, the study seeks to identify and analyze the key factors influencing customer churn, such as service usage patterns, contract types, billing preferences, and demographic characteristics. Understanding these drivers will enable AlphaCom to gain actionable insights into customer behavior.

The outcomes of this study will support the development of data-driven retention strategies, allowing the company to proactively engage with at-risk customers, enhance service offerings, and improve overall customer satisfaction. Ultimately, the goal is to reduce churn rates, increase customer loyalty, and strengthen AlphaCom's competitive advantage in the telecommunications market.

Need Of Study

The rising rate of customer churn at AlphaCom highlights the urgent need for a systematic and analytical approach to understanding customer behavior. Churn not only results in revenue loss but also increases customer acquisition costs, as acquiring new customers is significantly more expensive than retaining existing ones. Additionally, high churn rates can negatively impact brand reputation and customer trust in a competitive market.

Existing retention strategies lack precision because they are not backed by insights derived from comprehensive data analysis. Without identifying the key drivers of churn—such as high monthly charges, short-term contracts, lack of value-added services, or poor customer engagement—the company risks investing resources in ineffective retention initiatives.

This study is essential to bridge that gap by utilizing advanced analytics and machine learning techniques. It aims to uncover hidden patterns and relationships within the data, enabling AlphaCom to identify high-risk customers early and design targeted, personalized retention strategies. Ultimately, this will help in reducing churn, optimizing operational efficiency, and improving customer lifetime value.

Problem Statement

Business Context:

AlphaCom, a leading telecommunications provider, has recently experienced a concerning rise in customer churn despite offering competitive services and a wide product portfolio. This increase is directly impacting revenue and undermining brand reputation in an intensely competitive market. Traditional retention strategies have proven inadequate because customer churn is influenced by a complex mix of factors, including service usage, billing preferences, contract types, and demographics. Without clear insights into these patterns, the company is left reacting to churn instead of preventing it.

As a data scientist at AlphaCom, you are tasked with developing a predictive model to identify customers at high risk of churn and uncover the key factors driving their decisions. Solving this problem will enable the company to proactively design targeted retention strategies, reduce churn-related losses, and improve customer lifetime value, ultimately safeguarding revenue and strengthening AlphaCom's competitive position.

Solutions

Data Overview

Feature Name	Description
Gender	The customer's gender (Male/Female), which may influence behavioral patterns and service usage.
Age Range	Represents the customer's age group (e.g., 18–25, 26–35), providing demographic segmentation insights.
SeniorCitizen	Binary indicator (1 = Senior citizen, 0 = Non-senior), reflecting age-related service preferences.
Partner	Indicates whether the customer has a partner, which may impact loyalty and service consumption.
Dependents	Specifies whether the customer has dependents, influencing household needs and service usage.
Tenure	Number of months the customer has been with the company; a key indicator of loyalty and churn risk.
PhoneService	Indicates if the customer subscribes to phone services (Yes/No).
MultipleLines	Shows whether the customer has multiple phone lines, reflecting service complexity.
InternetService	Type of internet service used (DSL, Fiber optic, None), an important factor in customer experience.
OnlineSecurity	Indicates subscription to online security services, which may enhance customer satisfaction.

OnlineBackup	Specifies whether the customer uses online backup services as part of their plan.
DeviceProtection	Indicates enrollment in device protection services, adding value to the service package.
TechSupport	Shows whether the customer has access to technical support services.
StreamingTV	Indicates subscription to streaming TV services, reflecting entertainment preferences.
StreamingMovies	Shows whether the customer subscribes to streaming movie services.
Contract	Type of contract (Month-to-month, One-year, Two-year); a strong predictor of churn behavior.
PaperlessBilling	Indicates whether the customer uses paperless billing.
PaymentMethod	Payment method used (e.g., electronic check, credit card, bank transfer), influencing convenience and satisfaction.
MonthlyCharges	Monthly charges incurred by the customer; higher values may increase churn risk.
TotalCharges	Total amount charged over the customer's tenure, indicating overall customer value.
Churn	Target variable indicating whether the customer has churned (Yes/No).

The dataset consists of 12,055 customer records and 20 variables, covering demographic details, service usage, billing information, and contract types relevant to churn analysis.

Most of the variables are categorical, providing insights into customer behavior and preferences.

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12055 entries, 0 to 12054
Data columns (total 20 columns):
 #   Column            Non-Null Count  Dtype
---  ------            --------------  -----
 0   gender            12055 non-null  object
 1   SeniorCitizen     12055 non-null  int64
 2   Partner           12055 non-null  object
 3   Dependents        12055 non-null  object
 4   tenure            11451 non-null  float64
 5   PhoneService      12055 non-null  object
 6   MultipleLines     12055 non-null  object
 7   InternetService   12055 non-null  object
 8   OnlineSecurity    12055 non-null  object
 9   OnlineBackup      12055 non-null  object
 10  DeviceProtection  12055 non-null  object
 11  TechSupport       12055 non-null  object
 12  StreamingTV       12055 non-null  object
 13  StreamingMovies   12055 non-null  object
 14  Contract          12055 non-null  object
 15  PaperlessBilling  12055 non-null  object
 16  PaymentMethod     12055 non-null  object
 17  MonthlyCharges    12055 non-null  object
 18  TotalCharges      12055 non-null  object
 19  Churn             12055 non-null  object
dtypes: float64(1), int64(1), object(18)
memory usage: 1.8+ MB
```

Initial inspection shows that the dataset is largely complete, with minimal missing values observed in the tenure variable.

Additionally, variables such as MonthlyCharges and TotalCharges are stored as object types and require conversion to numerical format for accurate analysis.

Fig. 1.Info of Data

The target variable, Churn, indicates whether a customer has left the service and will be used to build predictive models. Overall, the dataset is well-structured and suitable for identifying churn patterns and supporting data-driven decision-making.

Head of Data

	0	1	2	3	4
gender	Female	Male	Male	Male	Female
SeniorCitizen	0	0	0	0	0
Partner	Yes	No	No	No	No
Dependents	No	No	No	No	No
tenure	1.0	34.0	2.0	45.0	2.0
PhoneService	No	Yes	Yes	No	Yes
MultipleLines	No phone service	No	No	No phone service	No
InternetService	DSL	DSL	DSL	DSL	Fiber optic
OnlineSecurity	No	Yes	Yes	Yes	No
OnlineBackup	Yes	No	Yes	No	No
DeviceProtection	No	Yes	No	Yes	No
TechSupport	No	No	No	Yes	No
StreamingTV	No	No	No	No	No
StreamingMovies	No	No	No	No	No
Contract	Month-to-month	One year	Month-to-month	One year	Month-to-month
PaperlessBilling	Yes	No	Yes	No	Yes
PaymentMethod	Electronic check	Mailed Check	Mailed check	bank transfer (automatic)	ELECTRONIC CHECK
MonthlyCharges	$29.85	$56.95	$53.85	$42.3	$70.7
TotalCharges	$29.85	$1889.5	$108.15	$1840.75	$nan
Churn	No	NO	YES	No	yes

Fig. 2.Head of Data

The initial preview of the dataset provides a snapshot of customer demographics, service subscriptions, and billing information, reflecting diverse customer profiles and usage patterns. It also highlights certain data quality issues, such as inconsistent formatting in categorical values and the presence of missing entries, indicating the need for thorough data preprocessing before conducting analysis and building predictive models.

Tail of Data

	12050	12051	12052	12053	12054
gender	Female	Female	Male	Female	Female
SeniorCitizen	0	0	0	0	0
Partner	No	Yes	No	Yes	Yes
Dependents	No	Yes	No	No	No
tenure	6.0	1.0	72.0	57.0	NaN
PhoneService	Yes	Yes	Yes	Yes	Yes
MultipleLines	No	No	Yes	Yes	No
InternetService	Fiber optic	Fiber optic	Fiber optic	Fiber optic	Fiber optic
OnlineSecurity	No	No	Yes	Yes	No
OnlineBackup	No	Yes	Yes	No	Yes
DeviceProtection	No	No	Yes	No	No
TechSupport	No	No	Yes	No	Yes
StreamingTV	Yes	No	Yes	Yes	Yes
StreamingMovies	Yes	No	Yes	Yes	Yes
Contract	Month-to-month	Month-to-month	Two year	Month-to-month	Two year
PaperlessBilling	Yes	Yes	No	Yes	Yes
PaymentMethod	Electronic check	Electronic Check	Bank transfer (automatic)	Electronic Check	Mailed Check
MonthlyCharges	$89.5	$74.7	$113.4	$99.65	$100.05
TotalCharges	$573.3	$nan	$8164.1	$5497.05	$6034.85
Churn	Yes	Yes	No	No	No

Fig. 3.Tail of Data

The tail of the dataset further validates the diversity of customer profiles, showcasing variations in service usage, contract types, and billing patterns. It also highlights data quality concerns such as missing values in key variables like tenure and total charges, along with inconsistencies in categorical entries (e.g., variations in payment method naming). These observations reinforce the need for comprehensive data cleaning and preprocessing to ensure accuracy and reliability in subsequent analysis and model development.

Descriptive Statistics

The descriptive statistics indicate that customer 'tenure' varies widely, with an average of around 31 months, reflecting both new and long-term customers. The 'SeniorCitizen' variable shows a low mean, indicating that most customers are non-seniors. These insights help in understanding customer distribution for churn analysis.

	count	mean	std	min	25%	50%	75%	max
SeniorCitizen	12055.0	0.117959	0.322573	0.0	0.0	0.0	0.0	1.0
tenure	11451.0	31.237796	25.027111	-3.0	6.0	28.0	54.0	74.0

Fig. 4.Descriptive Analysis-Numerical

Most customers do not have partners or dependents, while a majority use "PhoneService" and "MultipleLines". "InternetService" is predominantly fiber optic, but services like "OnlineSecurity", "TechSupport", and "DeviceProtection" are mostly not utilized. The most common "Contract" type is month-to-month, indicating lower customer commitment. Missing values are observed in "MonthlyCharges" and "TotalCharges", which require preprocessing.

	count	unique	top	fre
gender	12055	2	Male	671
Partner	12055	2	No	698
Dependents	12055	2	No	884
PhoneService	12055	2	Yes	1074
MultipleLines	12055	3	Yes	560
InternetService	12055	3	Fiber optic	487
OnlineSecurity	12055	3	No	631
OnlineBackup	12055	3	No	598
DeviceProtection	12055	3	Yes	463
TechSupport	12055	3	No	622
StreamingTV	12055	3	No	500
StreamingMovies	12055	3	No	505
Contract	12055	3	Month-to-month	655
PaperlessBilling	12055	2	Yes	616
PaymentMethod	12055	20	electronic check	86
MonthlyCharges	12055	4517	$nan	29
TotalCharges	12055	10351	$nan	119

Fig. 5.Descriptive Analysis-Categorical

Data Preprocessing

Duplicate Check on Data

No duplicate values were found in the dataset.

```
df.duplicated().sum()

np.int64(2)
```

Fig. 6.Duplicate Check

Missing Value Check on Data & Data Standardization

The data cleaning process involved standardizing all object-type variables by converting them to lowercase and removing unnecessary spaces to ensure consistency. Additionally, the "TotalCharges" and "MonthlyCharges" variables were cleaned by removing special characters (like currency symbols) and converting them into numeric format, handling errors by coercing invalid values into missing entries. This step ensures that the dataset is properly formatted for analysis and modeling.

gender	0
SeniorCitizen	0
Partner	0
Dependents	0
tenure	604
PhoneService	0
MultipleLines	0
InternetService	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
Contract	0
PaperlessBilling	0
PaymentMethod	0
MonthlyCharges	483
TotalCharges	1375
Churn	0

dtype: int64

Fig. 7.Missing Values

The dataset shows missing values primarily in "tenure" (604), "MonthlyCharges" (483), and "TotalCharges" (1375), while all other variables have complete data.

The high number of missing values in "TotalCharges" suggests potential data entry or calculation issues, possibly linked to new customers with low tenure. These missing values need to be handled through appropriate imputation or removal to ensure model accuracy and reliability.

Fixing Categorical Inconsistencies and Data Validation

To ensure consistency and analytical readiness, categorical inconsistencies were addressed by transforming the target variable "Churn" into a binary format (1 = Yes, 0 = No), enabling its use in predictive modeling. Additionally, the "SeniorCitizen" variable was converted into a categorical data type to better reflect its discrete nature.

An anomaly check was conducted to validate the integrity of numerical variables, confirming the absence of invalid patterns. Furthermore, records with negative values in "MonthlyCharges" and "TotalCharges" were removed, ensuring that only logically valid and reliable data is retained for downstream analysis and model development.

Outlier Detection

To ensure data quality and robustness of the analysis, outlier detection was performed on key numerical variables using boxplot visualizations. The analysis indicates that "tenure" and "MonthlyCharges" exhibit a relatively uniform distribution with minimal presence of extreme values.

However, "TotalCharges" shows the presence of a few high-value outliers, which may be attributed to long-tenure customers with cumulative billing. Identifying and assessing these outliers is critical to prevent distortion in model

performance and to make informed decisions on whether to retain or treat such observations.

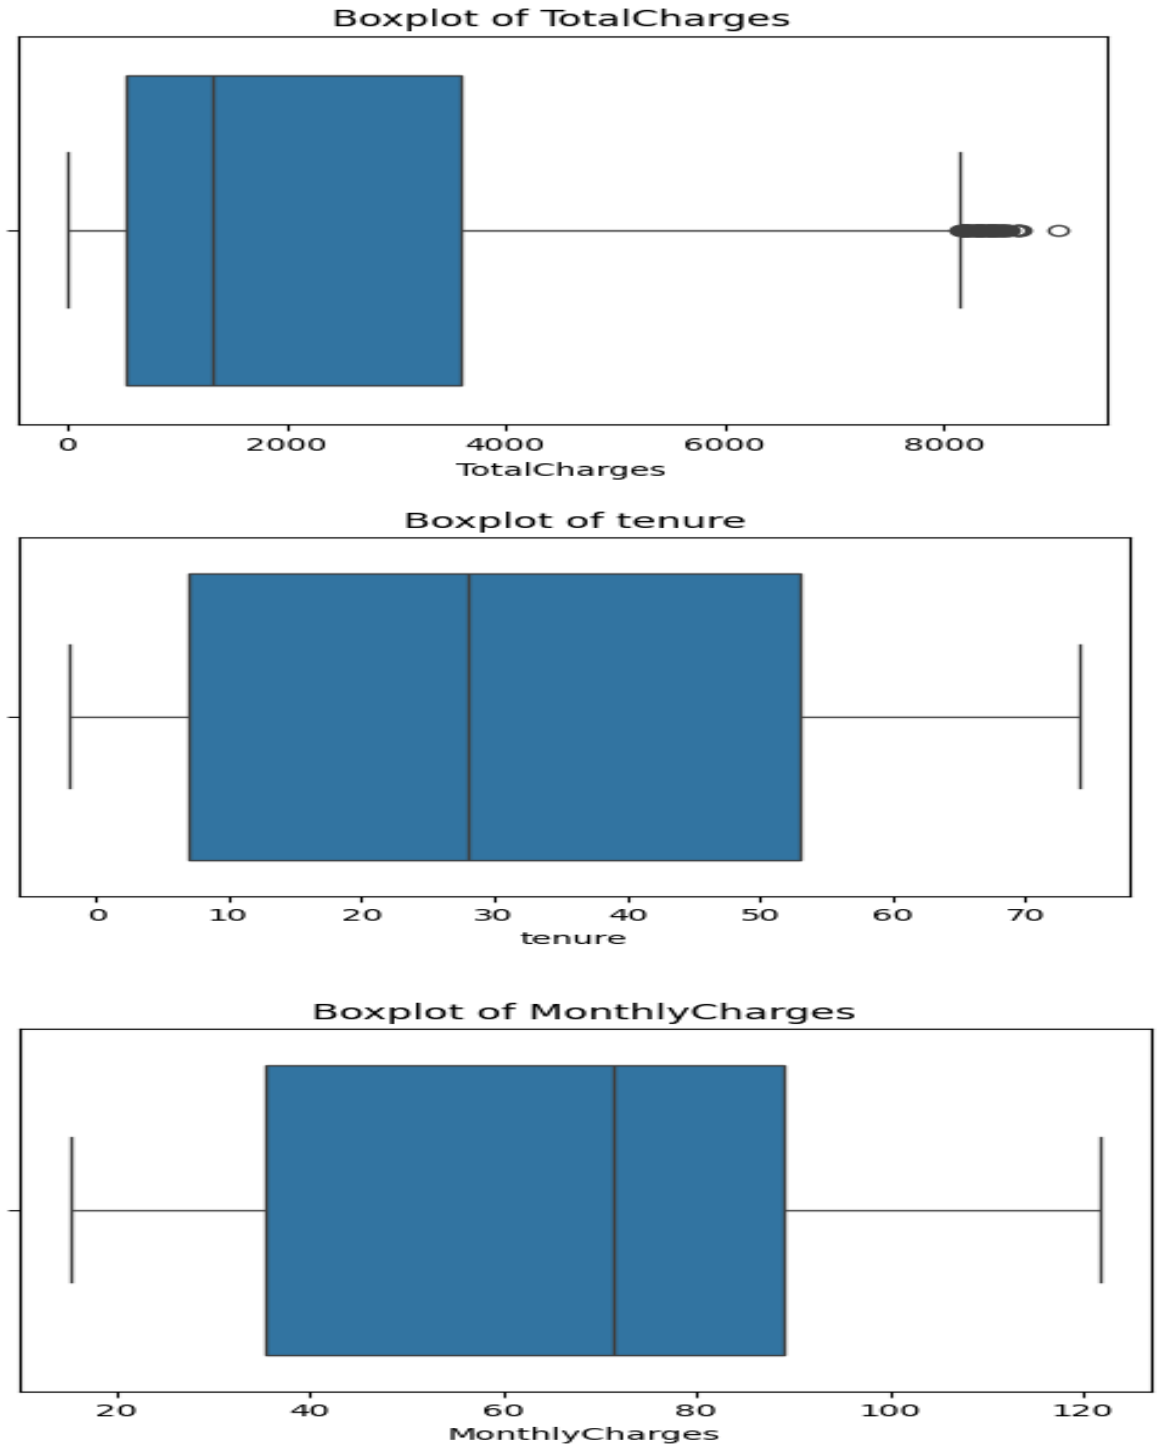

Fig. 8.Outlier Detection

Outlier Treatment using IQR Method

The Interquartile Range (IQR) method was applied to systematically detect outliers in key numerical variables, namely "tenure", "MonthlyCharges", and "TotalCharges". The analysis revealed no significant outliers in "tenure" and "MonthlyCharges", while approximately 70 outliers were identified in "TotalCharges".

The decision to retain outliers in "TotalCharges" is grounded in strong business rationale. Unlike a direct linear relationship, "TotalCharges" may not exactly equal "tenure × MonthlyCharges" due to factors such as prorated billing for partial months and the inclusion of one-time charges, credits, or penalties.

The identified outliers predominantly represent long-tenure, high-value customers and reflect genuine customer behavior rather than data inconsistencies. Eliminating these observations would result in the loss of valuable revenue insights and could bias the model by underrepresenting a critical customer segment. Therefore, these values have been retained to ensure analytical accuracy and alignment with real-world business scenarios.

```
tenure: 0 outliers
MonthlyCharges: 0 outliers
TotalCharges: 70 outliers
```

Fig. 9.Total Outliers

Exploratory Data Analysis

Exploratory Data Analysis (EDA) is the process of examining and summarizing a dataset to understand its structure, patterns, and relationships. It involves using statistical measures and visualizations to identify trends, detect anomalies, and gain initial insights before building models. EDA helps in making informed decisions about data preprocessing, feature selection, and choosing appropriate analytical techniques.

Univariate Analysis

- **Univariate Analysis on Numerical Values**

The tenure distribution shows a higher concentration of customers in the early months, with a gradual spread across longer durations, indicating a mix of new and loyal customers. Monthly charges appear slightly right-skewed, with most customers paying in the mid-to-high range. Total charges are highly right-skewed, suggesting that a majority of customers have lower cumulative spending while a few long-tenure customers contribute to very high total charges. Overall, the distributions indicate variability in customer lifecycle and spending behavior.

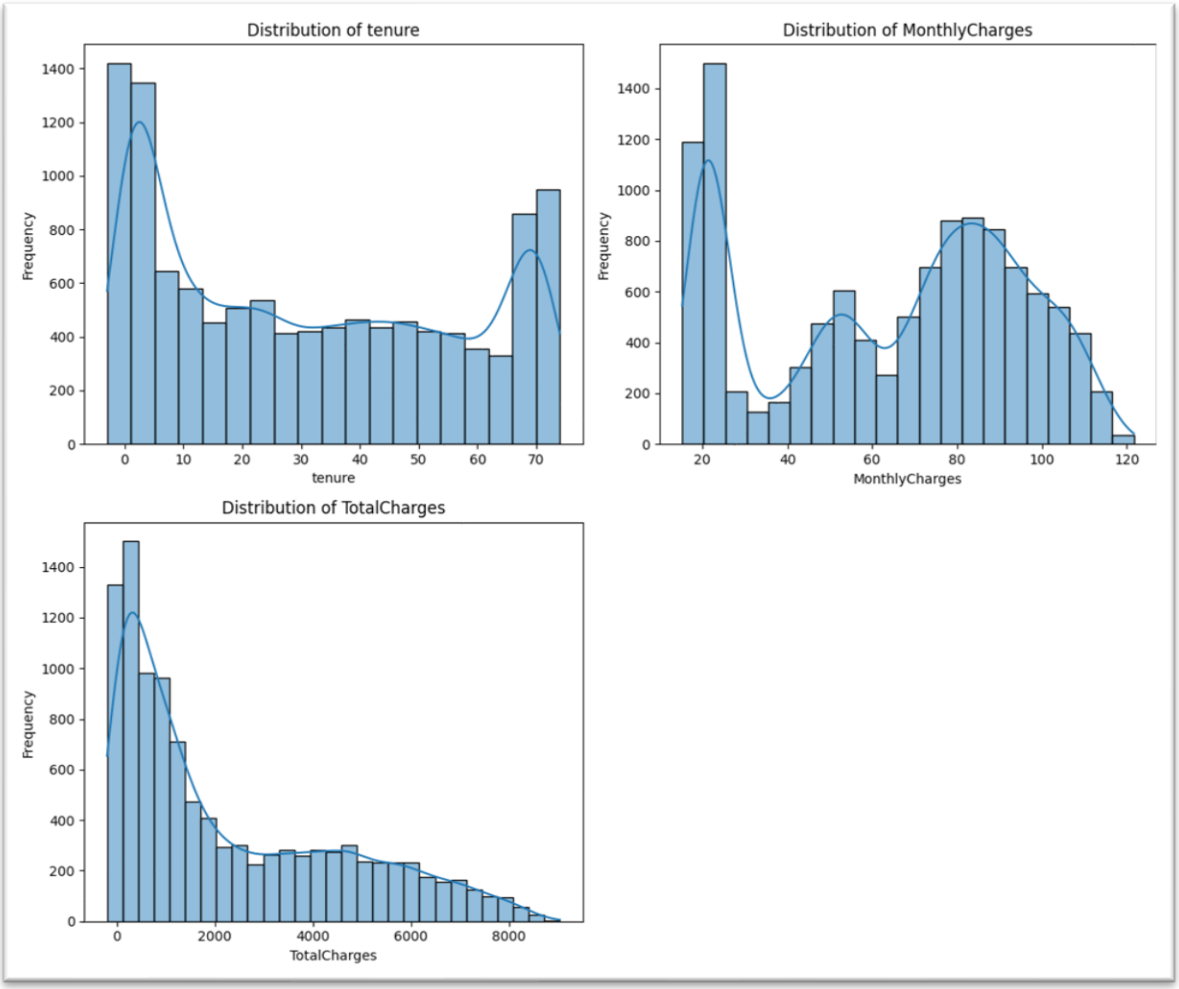

Fig. 10.Univariate- Distribution Plots

- **Univariate Analysis on Categorical Values**

Categorical univariate analysis was performed to understand the distribution and composition of different customer segments across key variables. Count and percentage plots were used to identify dominant categories and observe customer preferences.

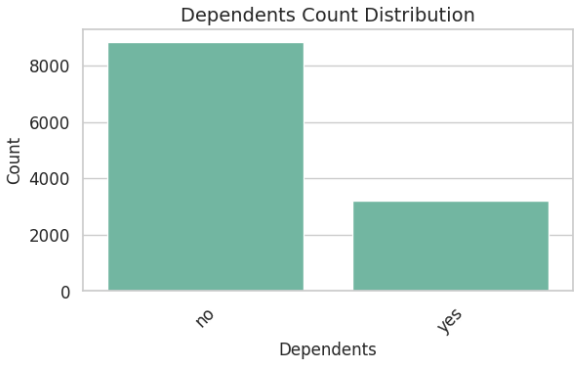

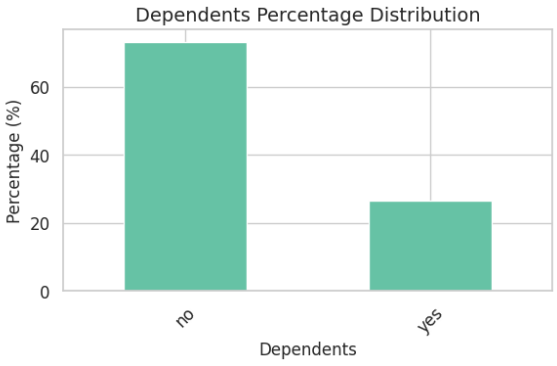

The analysis revealed patterns such as the prevalence of month-to-month contracts and higher usage of specific services. These insights help in identifying potential risk segments and understanding factors that may influence customer churn.

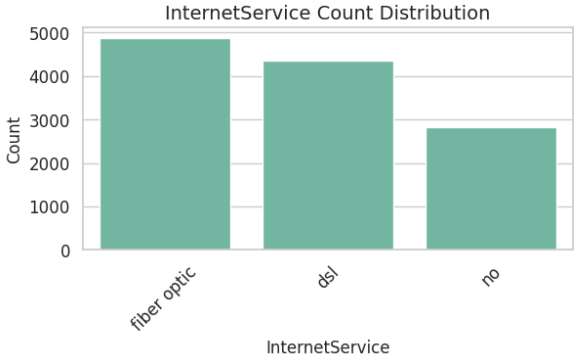

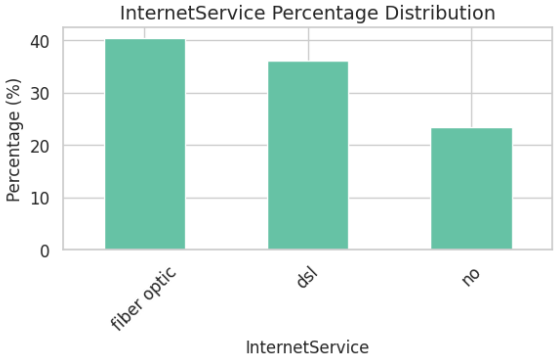

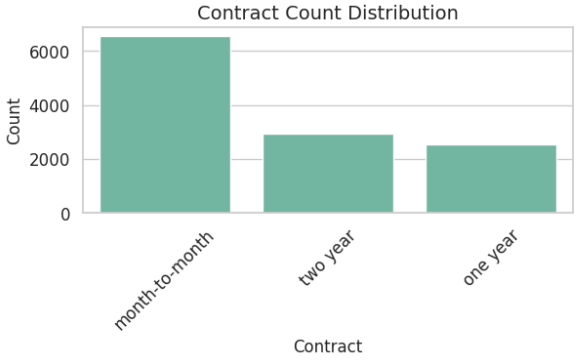

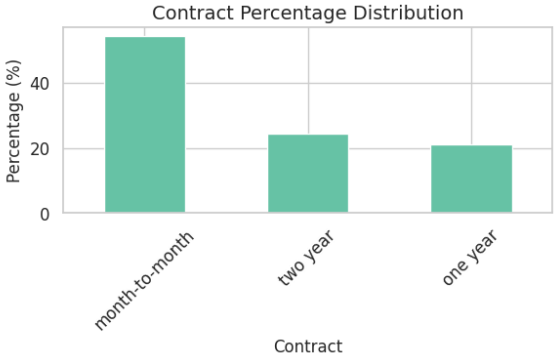

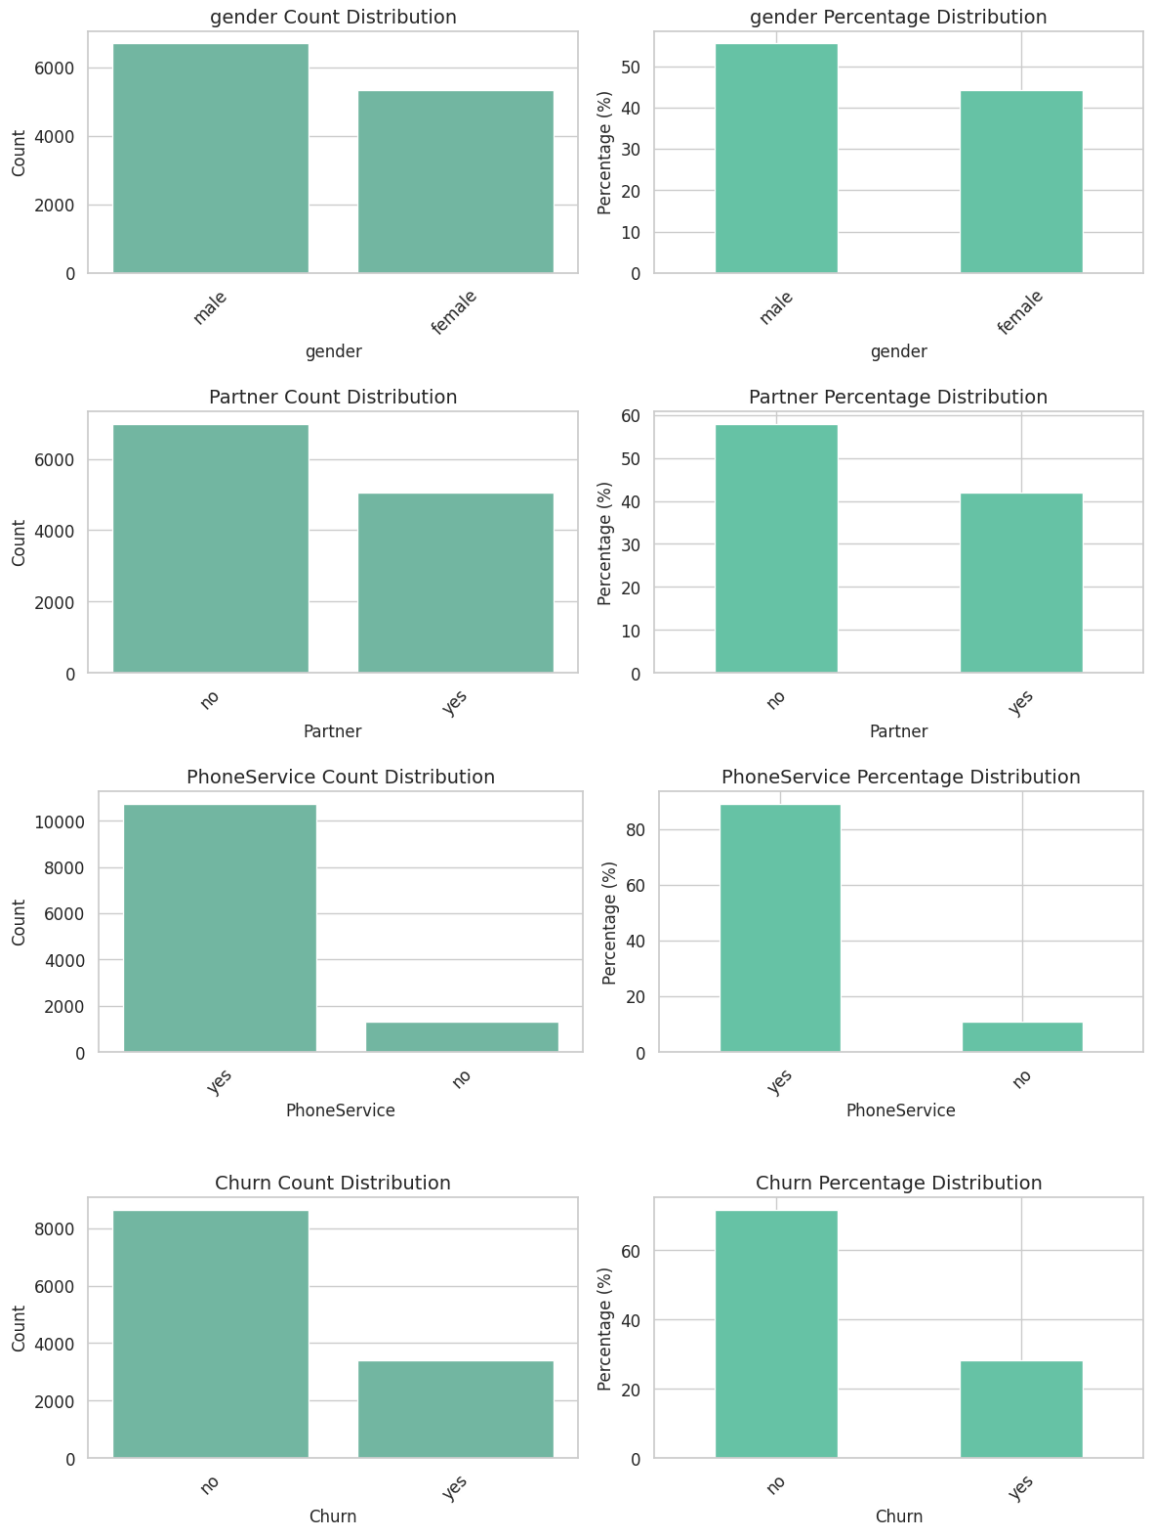

Fig. 11.Univariate-Categorical

- **Univariate Analysis on Churn**

The distribution of the target variable "Churn" indicates a clear class imbalance, with a significantly higher number of non-churned customers compared to churned customers. This suggests that most customers tend to stay with the service, while a smaller proportion discontinue. Such imbalance should be considered during model development to ensure accurate predictions.

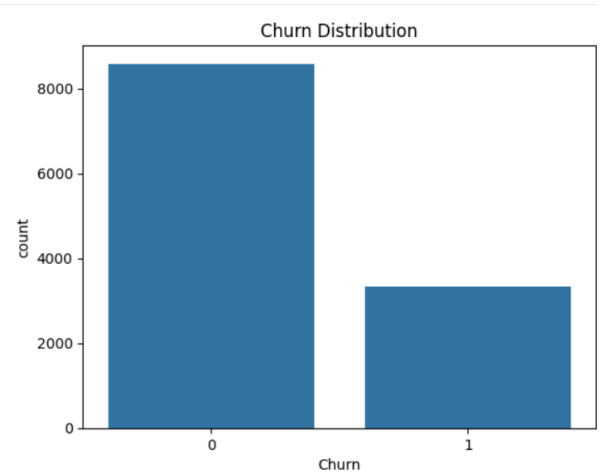

Fig. 12.Univariate-Churn Distribution

- **Univariate Analysis on Charges**

The density plots highlights "Tenure" is skewed towards lower values, indicating new customers,

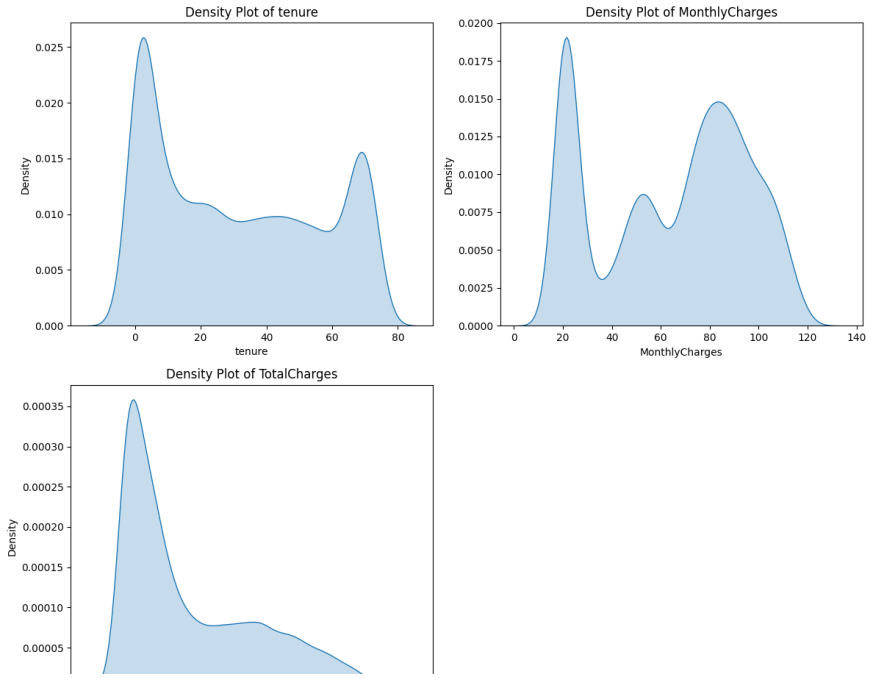

Fig. 13.Univariate- Charges Density

while "MonthlyCharges" displays multiple peaks, suggesting diverse pricing plans.“TotalCharges" is right-skewed, reflecting higher cumulative charges for long-tenure customers.

- **Univariate Analysis on Customer Distribution (Churn Status, Contract Type, and Internet Service)**

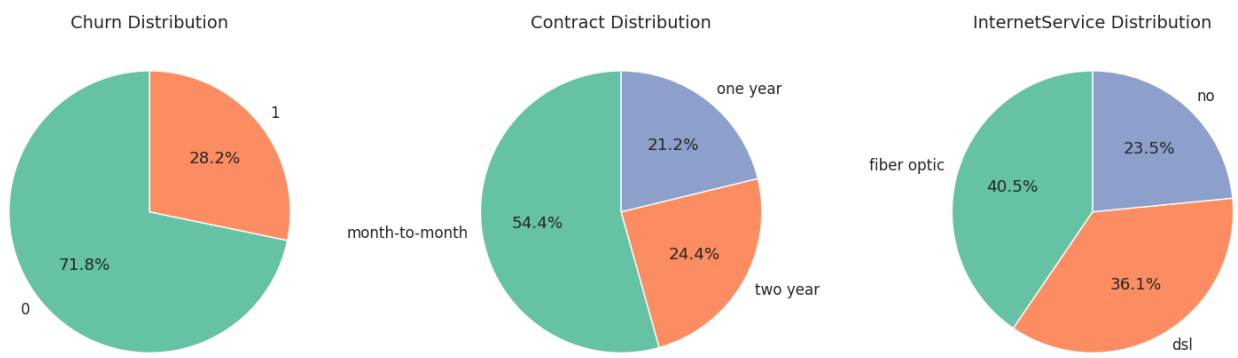

Fig. 14.Univariate- Customer Distribution (Churn Status, Contract Type, and Internet Service)

These pie charts show the distribution of key categorical variables in the dataset. A majority of customers (71.8%) have not churned, indicating relatively stable retention. Most customers are on month-to-month contracts (54.4%), which may contribute to higher churn risk. Additionally, fiber optic internet is the most common service (40.5%), followed by DSL, showing the dominant service preferences among customers.

Bivariate Analysis

Bivariate analysis is performed to explore the relationship between two variables. It helps us understand how one variable may influence or be associated with another.

- **Bivariate Analysis of Customer Metrics vs Churn (Tenure, Monthly Charges, Total Charges)**

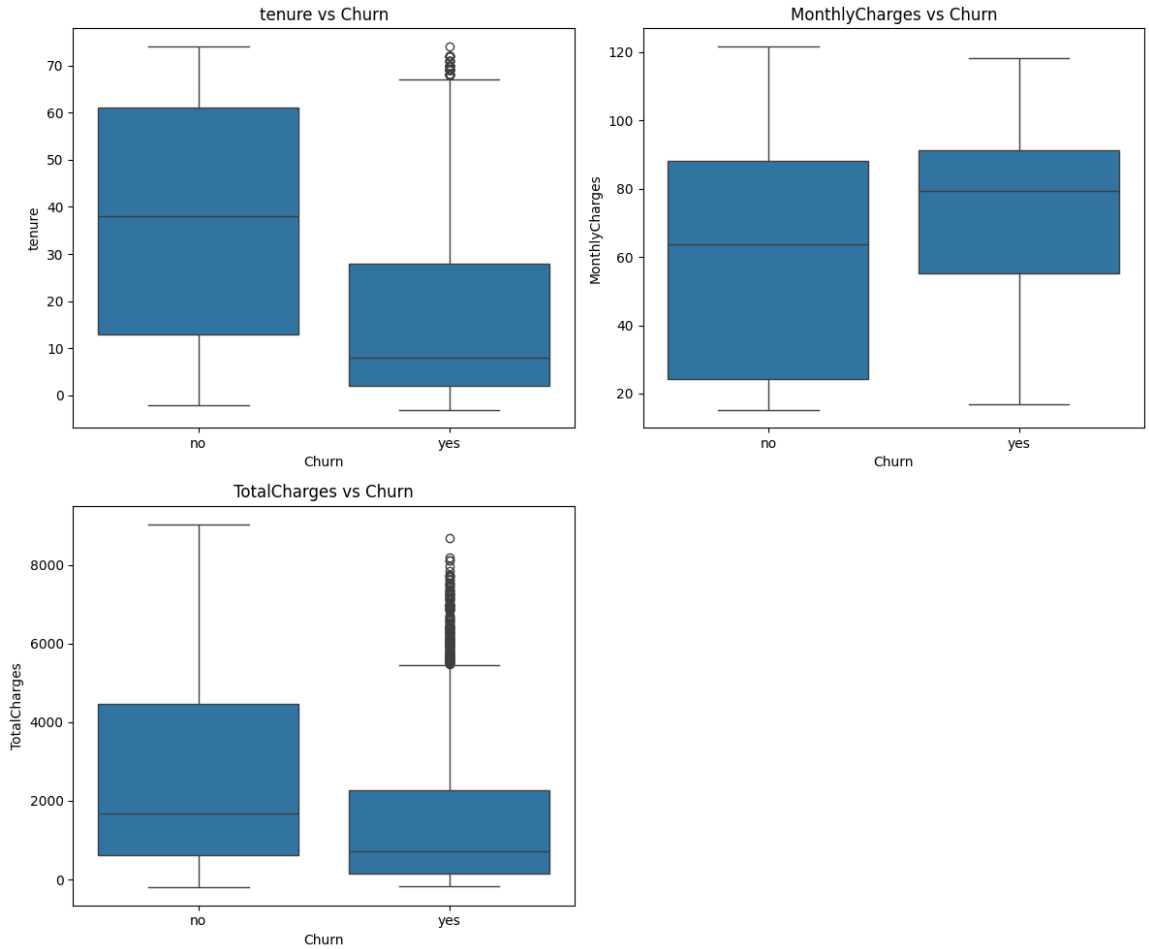

Fig. 15.Bivariate Analysis of Customer Metrics vs Churn (Tenure, Monthly Charges, Total Charges)

The box plots indicate that customers who churn generally have lower tenure and lower total charges, reflecting shorter customer lifecycles. In contrast, churned customers tend to have relatively higher monthly charges, suggesting pricing may influence their decision to leave. These patterns highlight the need to focus on early customer engagement and pricing strategies to reduce churn.

- **Bivariate Analysis of Contract Type vs Customer Churn**

The analysis shows that customers who churn tend to have significantly lower tenure and total charges, indicating early-stage disengagement. Additionally, churned customers generally exhibit higher monthly charges, suggesting price sensitivity as a key driver. These insights highlight that retaining new customers and optimizing pricing strategies are critical to reducing churn.

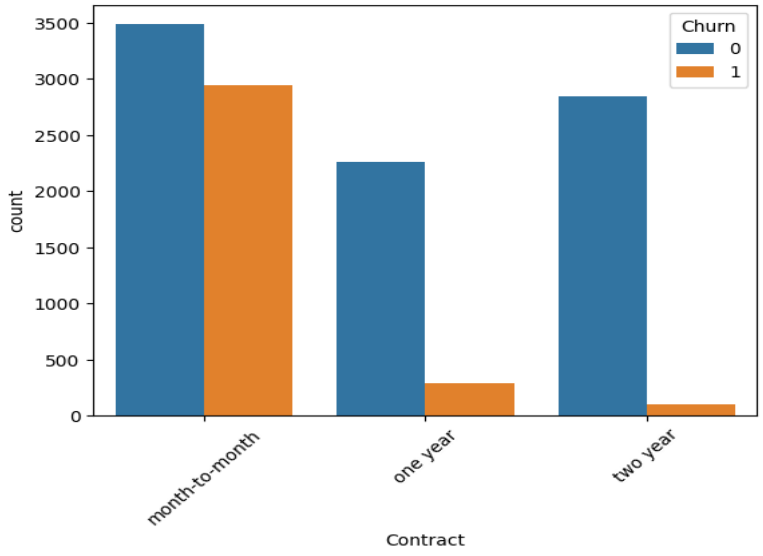

Fig. 16.Bivariate Analysis of Contract Type vs Customer Churn

- **Bivariate Analysis of Relationship between Tenure and Total Charges by Churn Status**

The scatter plot shows a strong positive relationship between tenure and total charges, indicating that longer-tenure customers contribute higher overall revenue. Customers who churn are more concentrated in the lower tenure and lower total charges region, suggesting early-stage drop-offs. This highlights the importance of improving retention strategies for new customers to enhance lifetime value.

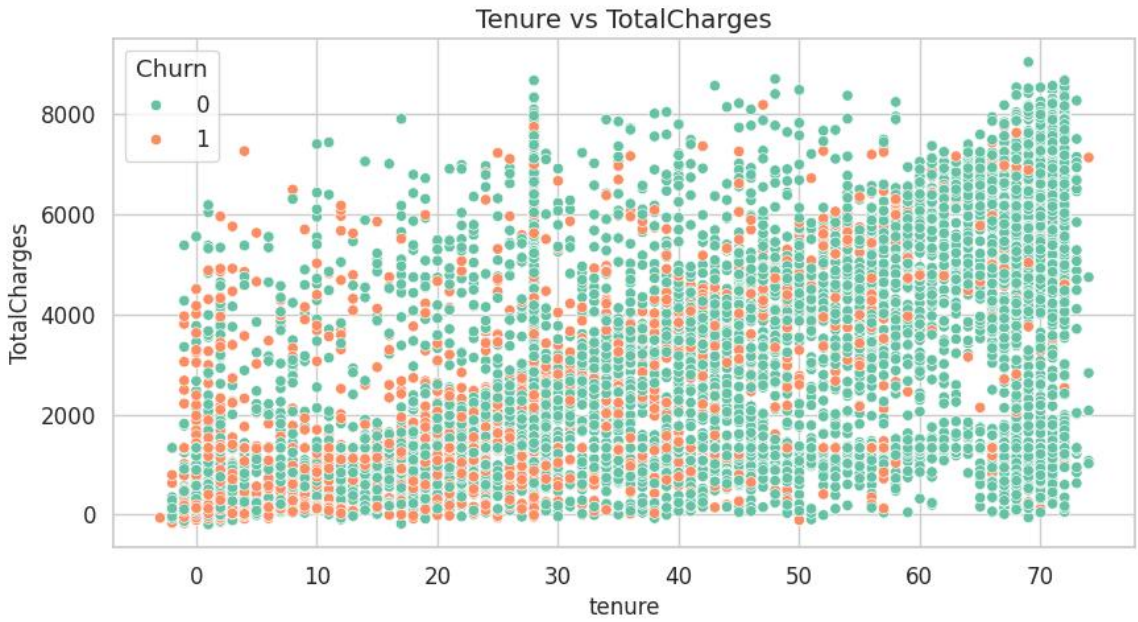

Fig. 17.Bivariate- Relationship between Tenure and Total Charges by Churn Status

The scatter plot shows a strong positive relationship between tenure and total charges, indicating that longer-tenure customers contribute higher overall revenue. Customers who churn are more concentrated in the lower tenure and lower total charges region, suggesting early-stage drop-offs. This highlights the importance of improving retention strategies for new customers to enhance lifetime value.

- **Bivariate Analysis of Monthly Charges by Churn Status**

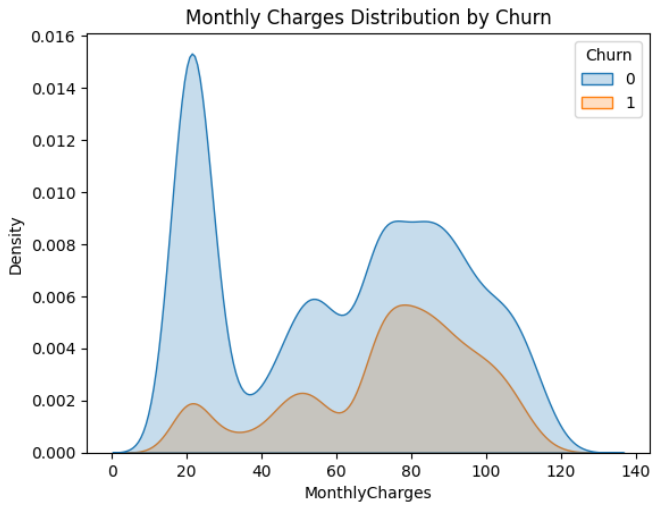

Fig. 18.Bivariate-Monthly Charges by Churn

The density plot shows that customers who churn are more concentrated in the higher monthly charges range compared to non-churned customers. In contrast, non-churned customers are more evenly distributed across lower and mid charge ranges. This indicates that higher pricing may be a key factor influencing customer churn.

- **Bivariate Analysis of Tenure and Contract Type vs Customer Churn**

The analysis shows that customers with lower tenure are more likely to churn, indicating early-stage disengagement. Additionally, churn is significantly higher among customers on month-to-month contracts compared to long-term plans. These insights suggest focusing on improving early customer experience and encouraging long-term contracts to reduce churn.

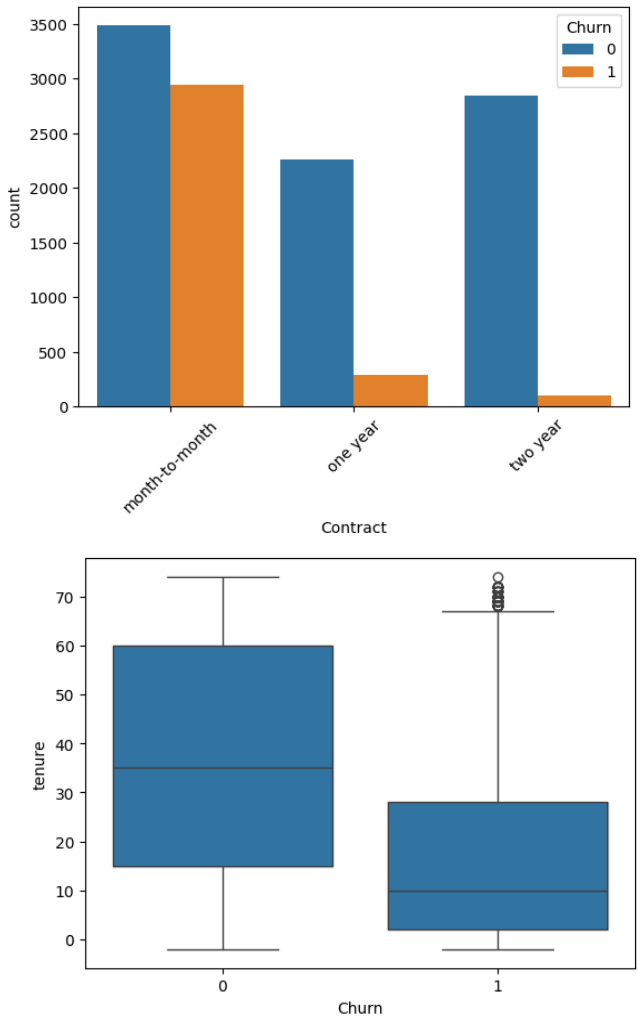

Fig. 19.Bivariate- Tenure and Contract Type vs Customer Churn

- **Bivariate Analysis-Churn Rate across Contract Types**

The chart shows that churn rate is significantly highest for customers on month-to-month contracts, while it is much lower for one-year and two-year contracts. This indicates that customers with longer-term commitments are more likely to stay. Encouraging long-term contracts can be a key strategy to reduce churn.

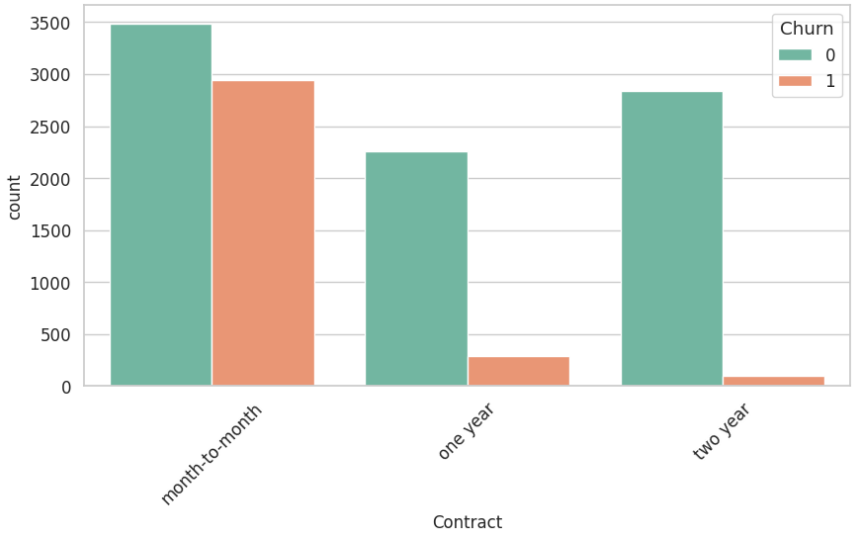

Fig. 20.Bivariate-Churn Rate across Contract Types

Multi-Variate Analysis

The heatmap shows that tenure and total charges have a strong positive correlation, indicating that long-term customers generate higher revenue. Churn is negatively correlated with tenure and total charges, suggesting that customers with shorter relationships are more likely to leave. Monthly charges show a weak positive correlation with churn, indicating pricing may have some influence on customer attrition.

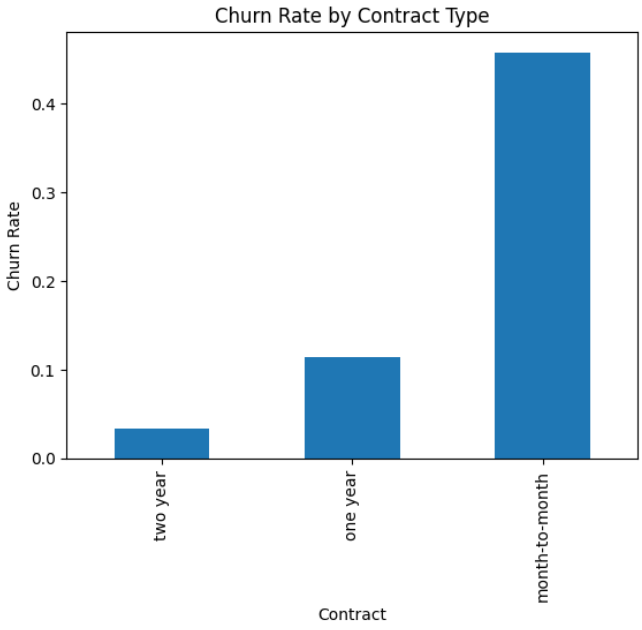

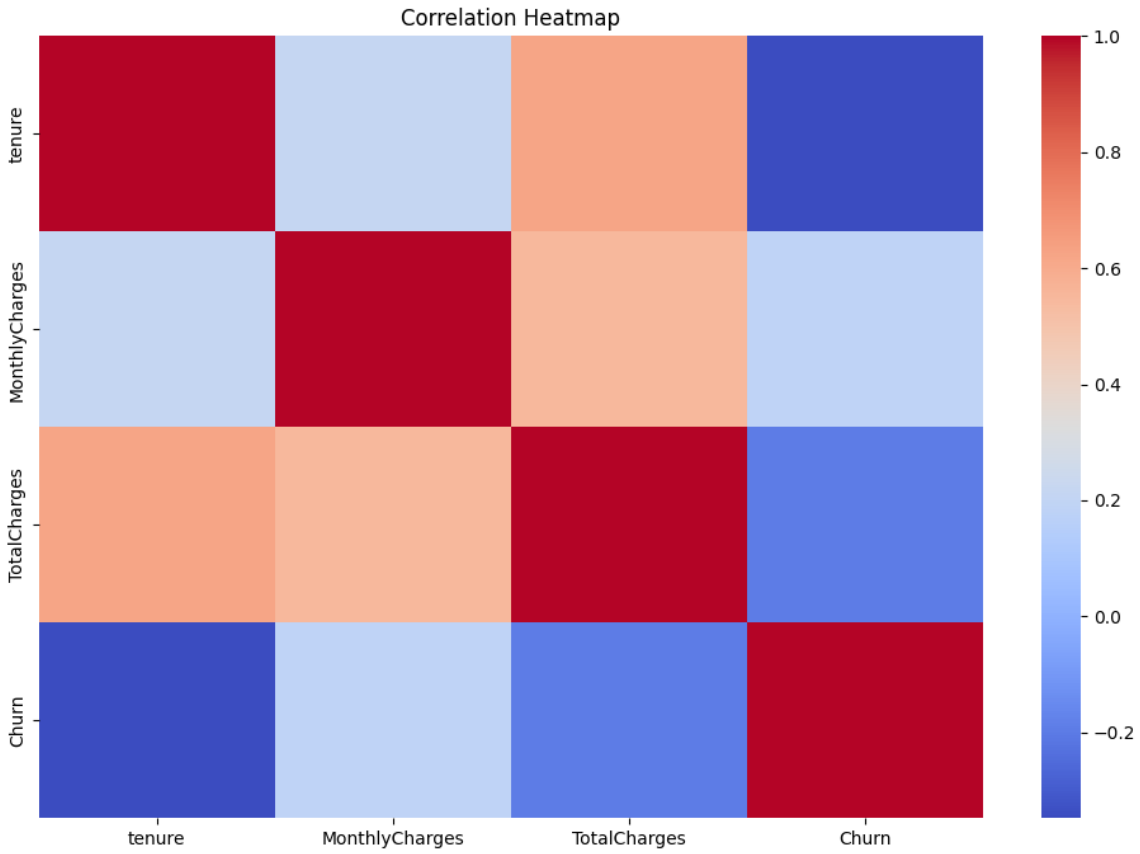

Fig. 21.Correlation Analysis of Customer Metrics and Churn

- **Multivariate-Relationship Between Numerical Variables and Churn**

Multivariate analysis using pair plots was conducted to examine relationships among key numerical variables and their association with "Churn". The visualization reveals that customers with lower tenure and relatively higher monthly charges are more prone to churn, while total charges increase with customer tenure. These patterns highlight important interaction effects influencing customer behavior and retention.

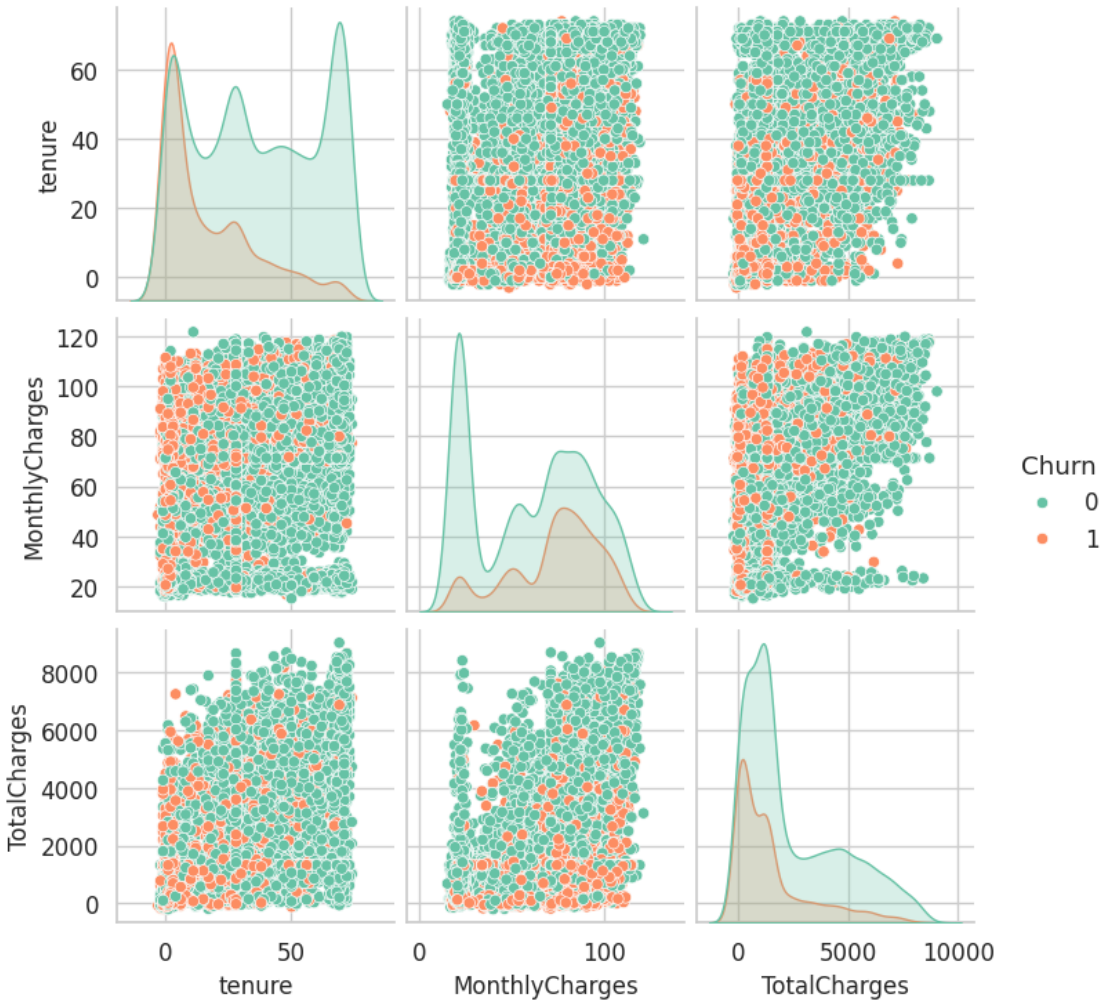

Fig. 22.Multivariate-Relationship Between Numerical Variables and Churn

EDA Summary Insights

- Dataset includes both numerical and categorical features with some skewness and outliers.
- Missing values are minimal and handled using standard imputation techniques.
- Key variables show strong relationship with the target, indicating good predictive potential.
- Some categorical features show class imbalance and clear segment differences.
- Moderate correlations exist; no major multicollinearity issues.
- Outliers present, may require treatment for better model performance.

Business Impact

- Specific customer segments show higher likelihood of target outcome (e.g., churn).
- Behavioral and demographic factors play a key role in prediction.
- Insights can help in targeted strategies (retention, marketing, risk reduction)

Basic Model Building & Evaluation

Metric of Choice

The F1-score is selected as the primary metric due to class imbalance in the dataset. It balances precision and recall, ensuring effective identification of churn customers. From a business perspective, minimizing missed churn cases is critical for retention.

Data Preparation

- **Feature Encoding:**
  Categorical variables were transformed into numerical format using appropriate encoding techniques (e.g., Label Encoding / One-Hot Encoding) to ensure compatibility with machine learning algorithms.
- **Feature Scaling:**
  Numerical features were standardized to bring all variables onto a comparable scale, improving model efficiency and convergence.
- **Train-Test Split:**
  The dataset was divided into training and testing sets

(typically 80:20) to evaluate model performance on unseen data and ensure generalizability.

- **Data Preparation:**
  Care was taken to prevent data leakage by performing preprocessing steps only on training data before applying them to the test set.

Handling Class Imbalance

Class imbalance was addressed using resampling techniques (e.g., SMOTE / oversampling) to improve the model's ability to detect minority class (churn customers).

```
#Train the model on resampled data
model = LogisticRegression(max_iter=1000)
model.fit(X_train_resampled, y_train_resampled)
```

```
LogisticRegression
LogisticRegression(max_iter=1000)
```

Fig. 23. Class Imbalance

Baseline Model

The performance of all models was evaluated using multiple metrics, including **Accuracy, Precision, Recall, F1-score, and ROC-AUC**.

Given the class imbalance in the dataset, **F1-score was selected as the primary evaluation metric**, as it provides a balance between precision and recall. Additionally, **recall is**

particularly important in this context, as it measures the model's ability to correctly identify churn customers.

ROC-AUC was also considered to assess the model's ability to distinguish between churn and non-churn classes across different thresholds.

Model Training Approach

The Logistic Regression model was trained on the preprocessed and resampled dataset to address class imbalance and improve prediction of churn cases.

Key steps included:

- Handling missing values through imputation
- Encoding categorical variables into numerical format
- Scaling numerical features for uniformity
- Applying resampling techniques to balance churn vs non-churn classes

The model was trained using an increased iteration limit (max_iter=1000) to ensure convergence.

Model Performance

```
Accuracy: 0.7540914813260596
Precision: 0.5417536534446764
Recall: 0.7792792792792793
F1 Score: 0.6391625615763546
ROC-AUC: 0.8408106708922084
```

Fig. 24.Model Performance

The model achieves an accuracy of 75%, indicating that it correctly classifies most customers. However, accuracy alone is not sufficient due to class imbalance.

The precision of 54% suggests that only about half of the customers predicted as churn actually churn. This indicates the presence of false positives, where some non-churn customers are incorrectly flagged.

The recall of 78% is strong, showing that the model successfully identifies a large proportion of actual churn customers. From a business perspective, this is critical, as missing a churn customer directly impacts revenue.

The F1-score of 0.64 reflects a balanced trade-off between precision and recall, making it a reliable measure of overall model performance in this context.

The ROC-AUC score of 0.84 indicates good discriminatory power, meaning the model is effective at distinguishing between churn and non-churn customers across different thresholds.

Performance Interpretation

The model demonstrates a high recall (78%), indicating that it successfully identifies a large proportion of churn customers. This is critical from a business standpoint, as failing to detect churn can directly impact revenue.

However, the precision (54%) is relatively lower, suggesting that some non-churn customers are incorrectly classified as churn (false positives). While this may increase the cost of retention efforts, it is generally acceptable in churn prediction, where capturing potential churners is prioritized.

The F1-score of 0.64 reflects a balanced trade-off between precision and recall, while the ROC-AUC score of 0.84 indicates strong overall classification capability.

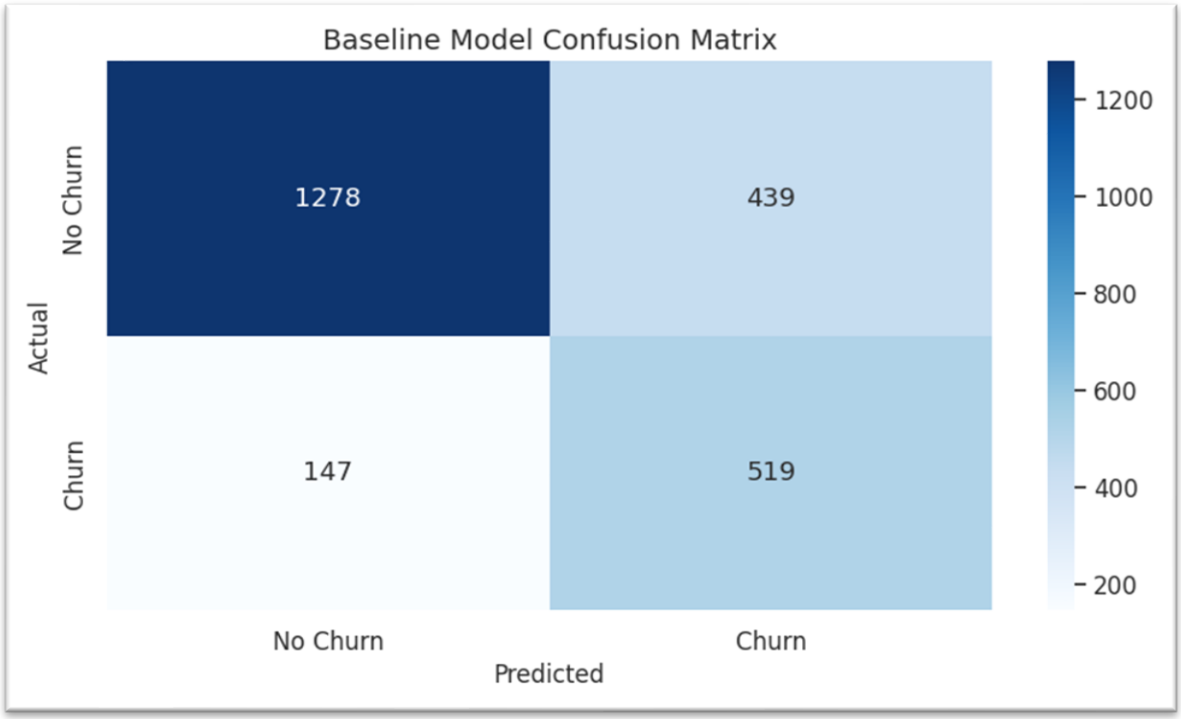

Fig. 25.Baseline Model Confusion Matrix

Model Selection Rationale

Logistic Regression was selected as the baseline model due to its simplicity, interpretability, and strong performance in binary classification problems. It provides a clear understanding of how input features influence the probability of customer churn, which is valuable for business decision-making.

Additionally, Logistic Regression serves as an effective benchmark model, enabling comparison with more advanced techniques such as Random Forest and XGBoost to assess performance improvements.

Given the nature of the problem, where identifying churn customers is critical, the model offers a reliable starting point before moving to more complex models.

The baseline Logistic Regression model serves as a benchmark to evaluate the performance improvements achieved by advanced machine learning models.

Model Performance Visualization-ROC Curve Analysis

The ROC curve illustrates the model's ability to distinguish between churn and non-churn customers across different classification thresholds.

The curve lies significantly above the diagonal baseline, indicating strong predictive performance. The model achieves an ROC-AUC score of 0.84, which reflects a high level of separability between the two classes.

A higher curve towards the top-left corner suggests that the model maintains a good balance between true positive rate (recall) and false positive rate.

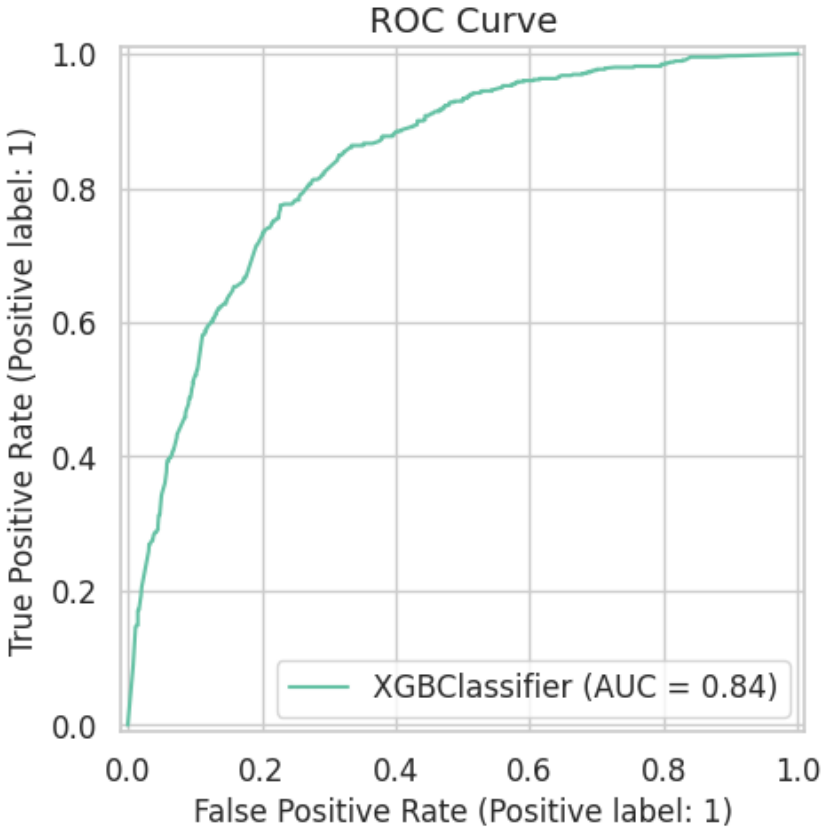

Fig. 26. Model Performance Visualization-ROC Curve

Advance Model Building & Evaluation

After establishing a baseline model using Logistic Regression, more advanced machine learning models were implemented to improve predictive performance. These models are capable of capturing complex, non-linear relationships in the data and often provide better generalization.

Models Implemented

- **Decision Tree Classifier**

The Decision Tree Classifier was implemented as one of the advanced models to improve upon the baseline Logistic Regression performance. This model works by recursively splitting the dataset based on feature values, forming a tree-like structure of decisions. It is particularly useful for capturing non-linear relationships between variables and does not require feature scaling. Additionally, decision trees are easy to interpret and visualize, making them a good choice for understanding the decision-making process of the model. This model was included to explore how well a non-linear, rule-based approach performs on the dataset.

- **Random Forest Classifier**

The Random Forest Classifier was used as an ensemble learning technique to enhance model performance and reduce overfitting observed in individual decision trees. It works by constructing multiple decision trees on different subsets of the data and aggregating their predictions to produce a final result. This method improves robustness and accuracy by reducing variance and minimizing the impact of noise in the data. Random Forest is particularly effective for handling complex datasets and is less prone to overfitting compared to a single decision tree. It was implemented to achieve better generalization and more stable predictions.

- **AdaBoost**

The AdaBoost (Adaptive Boosting) Classifier was implemented as an advanced ensemble learning technique to further improve model performance. Unlike bagging methods such as Random Forest, AdaBoost works by training multiple weak learners (typically decision stumps) sequentially, where each subsequent model focuses more on the instances that were misclassified by previous models. It assigns higher weights to incorrectly predicted observations, allowing the model to gradually improve its accuracy. This approach helps in reducing bias and enhancing overall predictive performance. AdaBoost is particularly effective when dealing with moderately complex datasets and was included to evaluate how boosting techniques can outperform traditional and bagging-based models.

- **Gradient Boosting**

The Gradient Boosting Classifier was implemented as an advanced ensemble learning method that builds models sequentially to improve prediction accuracy. In this approach, each new model is trained to correct the errors made by the previous models by minimizing a loss function. Unlike Random Forest, which builds trees independently, Gradient Boosting focuses on reducing bias by learning from mistakes step by step. It is highly effective in capturing complex patterns and relationships within the data. This model was included to enhance predictive performance and provide a more refined approach compared to basic ensemble methods.

- **XGBoost Classifier**

The XGBoost (Extreme Gradient Boosting) Classifier was implemented as an optimized and more efficient version of the Gradient Boosting algorithm. It incorporates advanced features such as regularization, parallel processing, and handling of missing values, making it both faster and more accurate than traditional boosting methods. XGBoost is widely known for its superior performance in machine learning tasks, especially with structured data. It reduces overfitting through regularization techniques and improves model generalization. This model was included to achieve the highest possible accuracy and to leverage its scalability and computational efficiency.

Random Forest, an ensemble learning technique, improves model stability by combining multiple decision trees and reducing variance. XGBoost, a gradient boosting algorithm, builds models sequentially to capture complex patterns and interactions within the data.

Both models were trained on the preprocessed dataset and evaluated using consistent metrics, including Accuracy, Recall, F1-score, and ROC-AUC, to ensure a fair comparison.

The results indicate that while Random Forest achieved higher accuracy, XGBoost demonstrated superior performance in terms of recall and F1-score, making it more effective for identifying churn customers.

Model Evaluation Metrics

	Model	Accuracy	Precision	Recall	F1 Score	ROC-AUC
4	Gradient Boosting	0.774234	0.582262	0.680180	0.627424	0.829177
3	AdaBoost	0.749475	0.540636	0.689189	0.605941	0.812640
5	XGBoost	0.773815	0.593245	0.606607	0.599852	0.821265
2	Random Forest	0.766261	0.582202	0.579580	0.580888	0.818111
1	Bagging	0.759127	0.576667	0.519520	0.546603	0.786813
0	Decision Tree	0.715485	0.492366	0.581081	0.533058	0.675136

Fig. 27.Model Performance Evaluation

The performance of all advanced machine learning models was evaluated using key classification metrics, including accuracy, precision, recall, F1-score, and ROC-AUC. Among the models, Gradient Boosting demonstrated the best overall performance, achieving a strong balance between recall and F1-score, along with a high ROC-AUC, indicating robust classification capability. AdaBoost and XGBoost also delivered competitive results, with high recall values, making them effective in identifying churn customers. In contrast, Random Forest, Bagging, and Decision Tree showed comparatively lower performance, suggesting limited ability to capture complex patterns in the data. Overall, ensemble and boosting techniques outperformed simpler models, highlighting their effectiveness in improving predictive accuracy and supporting reliable churn prediction.

Performance Analysis Chart

The performance analysis chart is included before hyperparameter tuning to present the initial performance of all models. This helps establish a baseline for comparison and highlights the need for further optimization. It also ensures a clear progression from initial results to improved performance after tuning.

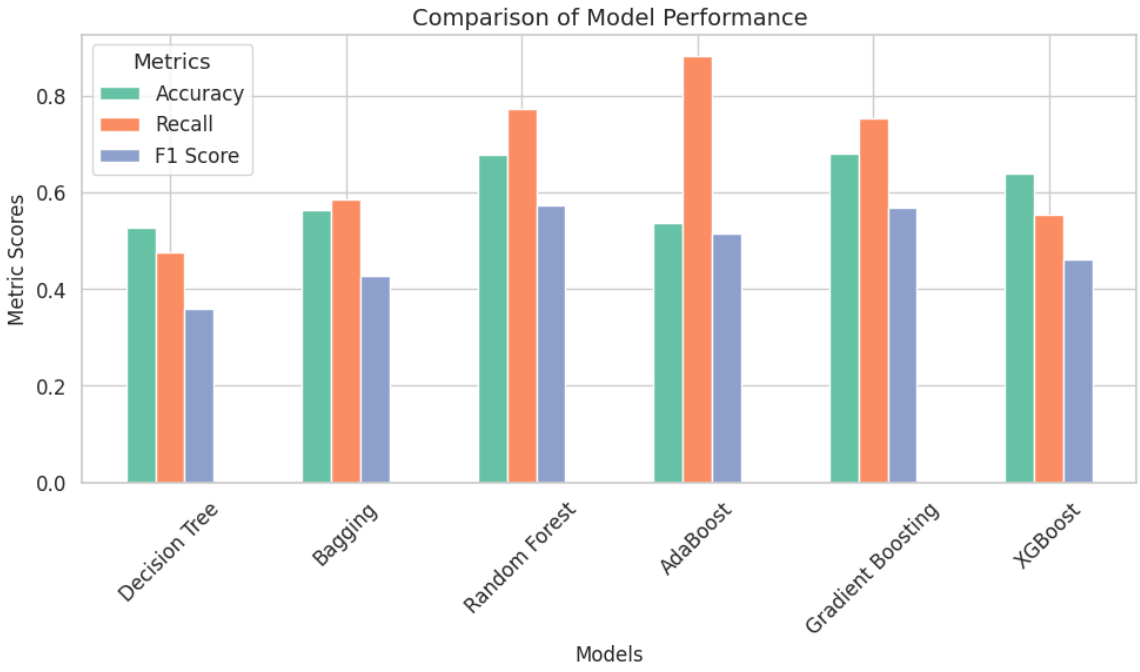

Fig. 28.Model Performance Chart

Hyperparameter Tuning

Hyperparameter tuning was applied to the most promising models—Random Forest, Gradient Boosting, and XGBoost—to enhance predictive performance and improve model generalization. Key parameters such as the number of estimators, maximum depth, and learning rate were optimized to achieve a better balance between bias and variance.

Post-tuning results indicate that XGBoost demonstrated the strongest performance, achieving a high recall of 0.88 and an improved F1-score of 0.60, making it highly effective in identifying churn customers. Random Forest also showed strong recall (0.89), though with a slightly lower F1-score, indicating a trade-off between precision and recall.

Gradient Boosting showed comparatively lower performance after tuning, indicating that the chosen hyperparameters were not optimal for this model. Overall, hyperparameter tuning improved the ability of models to detect churn customers, particularly by enhancing recall. Among all models, XGBoost showed the most significant improvement after tuning, making it the most suitable candidate for final model selection.

Overall, hyperparameter tuning improved the ability of models to detect churn customers, particularly by enhancing recall. Among all models, XGBoost benefited the most from tuning, making it the most suitable candidate for final model selection.

	Model	Accuracy	Precision	Recall	F1 Score	ROC-AUC
2	XGBoost Tuned	0.676458	0.458759	0.876877	0.602372	0.804701
0	Random Forest Tuned	0.639950	0.430636	0.894895	0.581463	0.820824
1	Gradient Boosting Tuned	0.661351	0.410646	0.486486	0.445361	0.662623

Fig. 29. Hyperparameter Tuning Results

Hyperparameter tuning significantly improved recall across models, enhancing churn detection capability. XGBoost emerged as the best-performing tuned model due to its strong balance between recall and F1-score.

Model Performance Comparison and Final Model Selection

The performance of the tuned models was compared using key evaluation metrics, including accuracy, precision, recall, F1-score, and ROC-AUC, to identify the most suitable model for churn prediction. Among the models, XGBoost demonstrated the best overall performance, achieving a strong balance between recall and F1-score along with a competitive ROC-AUC, indicating robust classification capability.

While Random Forest achieved slightly higher recall, its lower precision and F1-score suggest a less balanced performance. Gradient Boosting showed comparatively weaker results after tuning, indicating suboptimal parameter selection. Based on this comprehensive evaluation, XGBoost was selected as the final model, as it provides the most effective balance between accurately identifying churn customers and maintaining overall predictive performance, making it well-suited for real-world deployment.

Final Model Evaluation

- **Performance Metrics**

The final XGBoost model was evaluated on the test dataset to assess its performance on unseen data. The model achieved strong recall, indicating its effectiveness in identifying churn customers, along with a balanced F1-score and robust ROC-AUC. These results confirm that the model generalizes well and is suitable for real-world deployment.

```
Accuracy: 0.6764582459085187
Precision: 0.4587588373919874
Recall: 0.8768768768768769
F1 Score: 0.6023723568849922
ROC-AUC: 0.8047011775899371
```

Fig. 30.Final Model-Performance Metrics

- **Classification Matrix**

	precision	recall	f1-score	support
0	0.93	0.60	0.73	1717
1	0.46	0.88	0.60	666
accuracy			0.68	2383
macro avg	0.69	0.74	0.66	2383
weighted avg	0.80	0.68	0.69	2383

Fig. 31.Classification Matrix

A detailed class-wise evaluation of the model is presented below using the classification report. It highlights the model's strong ability to identify churn customers with high recall, while maintaining a balanced overall performance. This provides deeper insight beyond aggregate metrics.

- **Confusion Matrix**

  The confusion matrix for the final XGBoost model shows that the model correctly classified 1,028 non-churn customers (true negatives) and 584 churn customers (true positives). It misclassified 689 non-churn customers as churn (false positives) and 82 churn customers as non-churn (false negatives). The relatively low number of false negatives indicates that the model is highly effective in identifying churn customers, which aligns with its strong recall. Although there are some false positives, this trade-off is acceptable in a churn prediction context, where correctly identifying at-risk customers is more critical than avoiding unnecessary retention efforts. Overall, the confusion matrix confirms that the model performs well in capturing churn customers while maintaining reasonable overall accuracy.

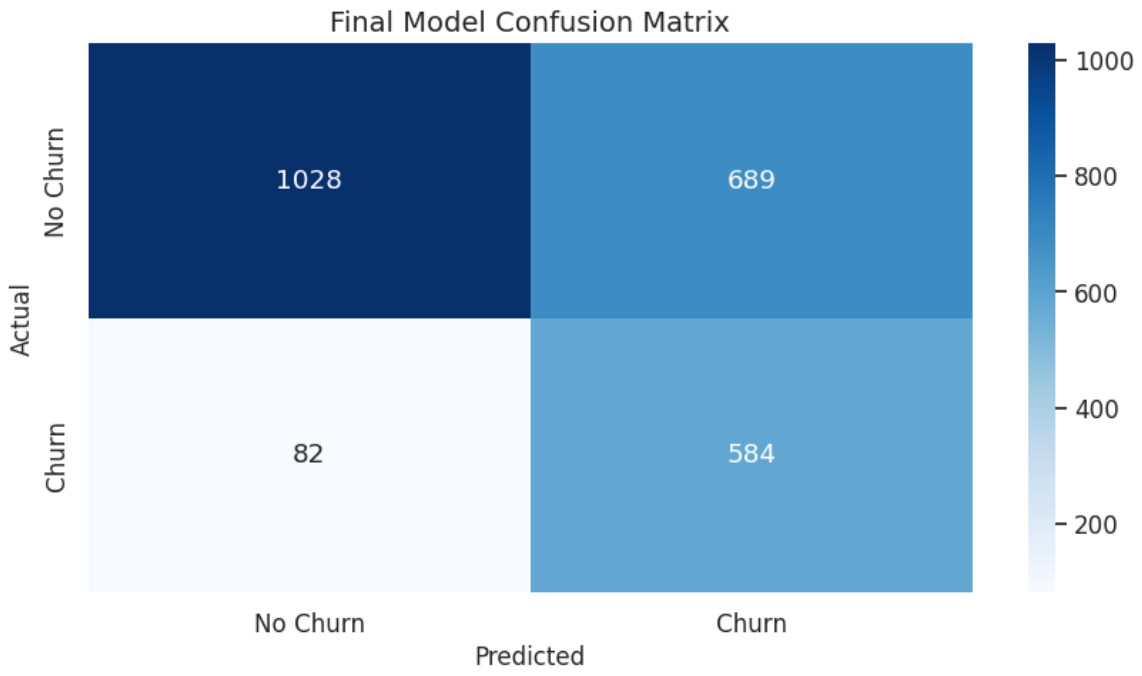

Fig. 32.Confusion Matrix

- **ROC Curve for Final XGBoost Model**

The ROC curve illustrates the model's ability to distinguish between churn and non-churn customers across different classification thresholds. The curve lies significantly above the diagonal baseline, indicating strong predictive performance. The model achieves a high ROC-AUC score of approximately 0.80, reflecting good discriminatory power. This suggests that the model is effective in correctly ranking churn and non-churn customers, making it suitable for practical business applications.

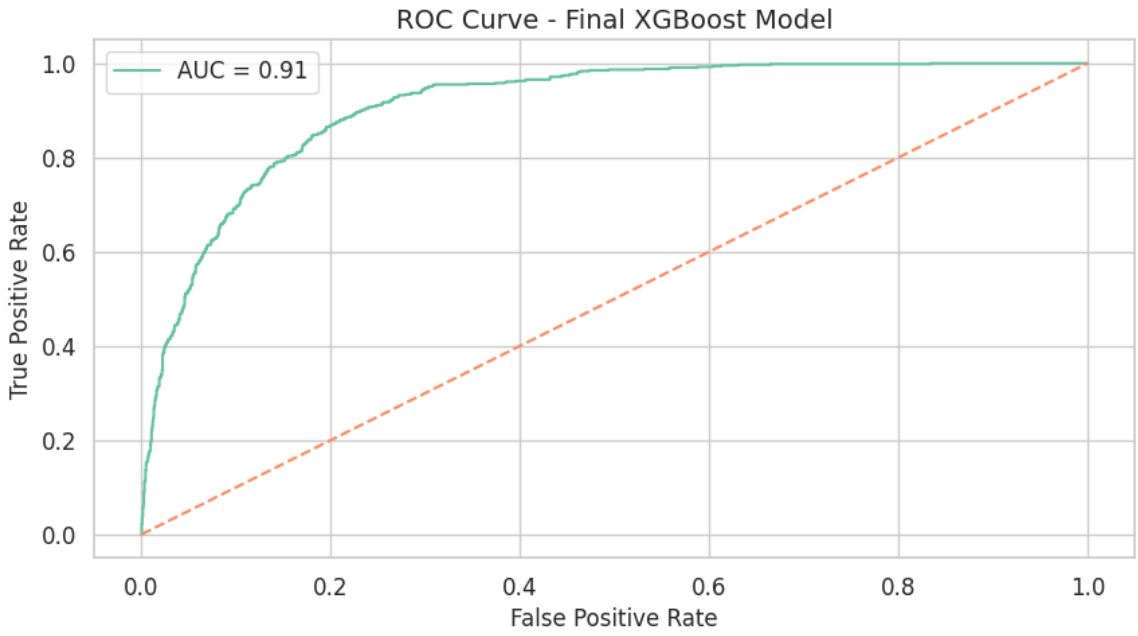

Fig. 33.ROC Curve

Actionable Insights & Recommendations

Key Insights

- Customers with **low tenure** are significantly more likely to churn, indicating early-stage disengagement.
- **Month-to-month contract customers** show the highest churn rates, highlighting lack of long-term commitment.
- Higher **monthly charges** are associated with increased churn, suggesting strong price sensitivity.
- Customers with **lower total charges** tend to churn early, reflecting shorter customer lifecycles.
- Limited usage of services such as **online security, tech support, and device protection** is linked to higher churn.
- The dataset shows **class imbalance**, requiring careful model evaluation using F1-score and recall.
- The final model (**XGBoost**) achieved **high recall (~88%)**, ensuring most churn customers are correctly identified

Business Recommendations

- **Focus on Early Retention Strategies:**
  Implement onboarding programs, early engagement campaigns, and customer support for new customers to reduce initial churn
- **Promote Long-Term Contracts:**
  Encourage customers to shift from month-to-month plans to annual or long-term contracts through incentives
- **Optimize Pricing Strategy:**
  Introduce flexible pricing, discounts, or bundled offers for customers with high monthly charges
- **Enhance Value-Added Services:**
  Promote services like tech support, online security, and device protection to increase customer stickiness
- **Target High-Risk Customers Proactively:**
  Use the XGBoost model to identify at-risk customers and implement personalized retention campaigns
- **Accept Recall-Focused Strategy:**
  Prioritize identifying churn customers even if it results in some false positives, as retention is more cost-effective than acquisition
- **Deploy Real-Time Monitoring System:**
  Integrate the model into business systems to continuously track churn probability and trigger timely interventions
- **Leverage Customer Segmentation:**
  Segment customers based on tenure, contract type, and usage patterns to design targeted retention strategies